

Relations - Querying



pm jat @ daiict



Relations - Querying

- Following are languages, that are available for querying relations
- In Theory
 - Relational Algebra
 - Relational Calculus
- In Practice
 - Standard Query Language (SQL)



Relational Algebra

- Queries are expressed as algebraic expression, where operands are relations.
- There is definite set of operators defined on relations. Some are unary while others are binary. For example sigma (σ) is an unary operation, and expressed as-

$\sigma_{DNO = 4 \text{ AND } SALARY > 50000}$ (employee)

- This expression evaluates to employee tuples that have value for DNO attribute=4, and SALARY > 50000
- Note: relation (relation instances) are operand to relational operators



Relational Algebra

- Relational Algebra operations are always “immutable”, in the sense that operands are not modified, even by unary operations, result is a new relation constructed by the operation
- Relational algebra expressions have closure property that is result of an expression is a valid relation.
- Relational algebra was introduced in original proposal of relational model.
- Complex queries are expressed by sequencing multiple operations.



Relational Calculus

- A query in relational calculus is expressed as a logic expression; example below expresses same query -

$\{t \mid \text{EMPLOYEE}(t) \text{ AND } t.\text{salary} > 50000 \text{ AND } t.\text{dno}=4) \}$

- Other is below that outputs specific attributes of tuple tuple t , where t is such that member of employee, and there $t.\text{dno}$ and $d.\text{dno}$ matches where d is tuple of department relation, and $d.\text{name}$ is 'Research'

$\{t.\text{Fname}, t.\text{Lname}, t.\text{Address} \mid \text{EMPLOYEE}(t) \text{ AND } (\exists d)(\text{DEPARTMENT}(d) \text{ AND } d.\text{Dname}=\text{'Research'} \text{ AND } d.\text{Dnumber}=t.\text{Dno}))\}$

- Note that calculus expressions are more declarative (rather than procedural) in nature; no concept of operation and their order.



Standard Query Language (SQL)

- SQL was not part of original proposal of Relational Model. Some people pronounce SQL as “SEQUEL”.
- Originally developed at IBM with System R. System R was first implementation of Relational Model in Early 70s.
- SQL has evolved since then, and now a standard language for relational databases.
- ANSI has published SQL-86, SQL-89, SQL-92, SQL-99, SQL-2003, SQL-2006, and SQL-2008, 2013 or so. SQL that discussed in most texts is SQL-99.
- Despite ANSI SQL standard, RDBMS often do some kind of deviation to it. Postgres’s SQL is probably most ANSI compliant.



Standard Query Language (SQL)

- SQL provides following categories of commands:
- Data Definition Language (DDL)
 - Used for defining database schema.
- Data Manipulation Language (DML)
 - Update commands
 - Querying commands
- Physical level commands
 - Used for improving performance of database.
- Transaction Control
 - Begin transaction, commit transaction, rollback etc
- Authorization
 - Grant permission to other users for access (read/update /delete, etc.)



Querying in this course

- Relational Algebra and SQL side by side
- Write SQL queries through Algebraic route
 - Find Algebraic solution and then
 - Translate into SQL
- Note: SQL was initially influenced by tuple relation calculus, however later updates incorporated most algebraic concepts as well.



Basic Querying operations

- Let us first discuss what is called as “the original operators” –
 - SELECT
 - PROJECT
 - Join
 - Cross JOIN
 - Set Operations: UNION, INTERSECTION, and DIFFERENCE
 - Aggregate operations
- Note: Update operations were not part of original relational model



SELECT operation

- SELECT is unary Operation; represented by sigma (σ)
- It is used to get subset of tuple from a relation for specified tuple selection criteria. Expressed as -

$$\sigma_{\langle \text{tuple selection criteria} \rangle}(R)$$

- For example: to select the EMPLOYEE tuples whose department number is four, we express as following-

$$\sigma_{DNO=4}(EMPLOYEE)$$

- Interpretation of above statement could be as following -
“produce a new relation by reading EMPLOYEE relation, and selecting only the tuple that meet the specified condition”



SELECT operation

- Same example with little extension in “tuple selection criteria”. Produce a relation from Employee relation by selecting tuples where DNO=4 and Salary > 30000

$$\sigma_{DNO=4 \text{ AND } SALARY>30000}(EMPLOYEE)$$

- The Resultant relation of SELECTION operation has-
 - Schema same as the operand relation
 - Degree is same as the operand relation
 - Cardinality can be anything from 0 to N. Where N is cardinality of operand relation.



PROJECT Operation

- PROJECT operation selects certain “columns” of a relation
- It can be viewed as getting vertical subset of a relation.
- Relational Algebra uses pi (π) operator for this operation.
- The general form of the PROJECT operation is

$$\pi_{\langle \text{attribute-list} \rangle}(R)$$

- Where $\langle \text{attribute list} \rangle$ is the desired list of attributes from the attributes of relation R.
- Example: $\pi_{fname,lname,salary}(EMPLOYEE)$



PROJECT Operation

- PROJECT is again a unary operator.
- Result of PROJECT operation –
 - is a valid relation – no duplicates
 - has cardinality same as of the instance of EMPLOYEE (and less, if projection operation has to drop some duplicate tuples in result set)
 - has degree equal to number of attributes specified in the list



More examples of Select and Project

$\sigma_{(DNO=4 \text{ AND } SALARY>30000) \text{ OR } (DNO=5 \text{ AND } SALARY > 30000)}$
(EMPLOYEE)

SELECT * FROM EMPLOYEE
WHERE (*DNO=4 AND SALARY>30000*)
OR (DNO=5 AND SALARY > 30000);

$\pi_{gender,salary}(EMPLOYEE)$
SELECT DISTICT GENDER, SALARY FROM EMPLOYEE;



Basic Querying operations

- Algebraic operations are immutable
- Result of any expression is a “valid relation”. Result has-
 - its own schema
 - its own instance
 - However NOT stored in database, just in working memory
 - Tuple in result relation are unique



SELECT and PROJECT operations in SQL



SELECT statement in SQL

- SELECT statement of SQL is not only SELECT operation of relational algebra.
- It is a SQL command for executing all “Query Answering” operations.
- Following is most basic form of SELECT statement of SQL

```
SELECT ____<attribute-list>____ FROM ____<relational-expr>____  
  
[WHERE ____<tuple selection criteria>____];
```

- Above expression is like a combination of SELECT and PROJECT of relational algebra; where you specify “tuple selection criteria”, and also specify “attribute list”.



SELECT and PROJECT operation SQL

SELECTION operation

RA: $\sigma_{DNO=4}(EMPLOYEE)$

SQL: **SELECT * FROM EMPLOYEE WHERE DNO = 4;**

Note ‘*’ for attribute list; it implies all attributes of operand relation, that is “relation” in FROM clause.

RA: $\sigma_{DNO=4 \text{ AND } SALARY > 30000}(EMPLOYEE)$

SQL: **SELECT * FROM EMPLOYEE WHERE DNO=4
and SALARY > 30000;**

- Note: following SQL statement all tuples of EMPLOYEE relation
SELECT * FROM EMPLOYEE;



SELECT and PROJECT operation SQL

PROJECT operation:

$$\pi_{fname, lname, salary}(EMPLOYEE)$$

is expressed in SQL as following-

SELECT fname, lname, salary FROM employee;

- Note: SQL may produce duplicate tuples in result relation? SQL require using “DISTINCT” for removing duplicate tuples. Observe difference between following two SQL statements-

SELECT dno FROM employee;

SELECT DISTINCT dno FROM employee;



SELECT and PROJECT operation SQL

SELECT and PROJECT together in SQL

SELECT fname, lname, salary FROM employee
WHERE DNO=4 AND SALARY > 30000;

- In algebra, it may typically be expressed as-

$$\pi_{fname, lname, salary}(\sigma_{DNO=4 \text{ AND } SALARY > 30000}(EMPLOYEE))$$

- Or, in two steps as following-

$$r1 \leftarrow \sigma_{DNO=4 \text{ AND } SALARY > 30000}(EMPLOYEE)$$

$$result \leftarrow \pi_{fname, lname, salary}(r1)$$



SELECT and PROJECT operation SQL

- SQL allows produce tuples in desired order, and provides ORDER BY clause in SELECT statement.
- Example below-

```
SELECT ssn, fname, lname, salary  
FROM employee  
WHERE dno = 5;  
ORDER BY SALARY DESC;
```

- Note: There is no ordering operation in Relational Algebra.



Properties of SELECT operation

- The SELECT operation $\sigma \langle \text{selection condition} \rangle (R)$ produces a relation $r(S)$ that has the same schema as R
- The SELECT operation σ is commutative; i.e.,
$$\sigma \langle \text{cond1} \rangle (\sigma \langle \text{cond2} \rangle (R)) = \sigma \langle \text{cond2} \rangle (\sigma \langle \text{cond1} \rangle (R))$$
- A cascaded SELECT operation may be applied in any order; i.e.,
$$\begin{aligned} &\sigma \langle \text{cond1} \rangle (\sigma \langle \text{cond2} \rangle (\sigma \langle \text{cond3} \rangle (R))) \\ &= \sigma \langle \text{cond2} \rangle (\sigma \langle \text{cond3} \rangle (\sigma \langle \text{cond1} \rangle (R))) \end{aligned}$$
- A cascaded SELECT operation may be replaced by a single selection with a conjunction of all the conditions; i.e.,
$$\begin{aligned} &\sigma \langle \text{cond1} \rangle (\sigma \langle \text{cond2} \rangle (\sigma \langle \text{cond3} \rangle (R))) \\ &= \sigma \langle \text{cond1} \rangle \text{ AND } \langle \text{cond2} \rangle \text{ AND } \langle \text{cond3} \rangle (R) \end{aligned}$$



Properties of PROJECT Operation

- The number of tuples in the result of projection $\pi_{\langle \text{list} \rangle} (R)$ is always less or equal to the number of tuples in R .
- If the list of attributes includes a key of R , then the number of tuples is equal to the number of tuples in R .
- $\pi_{\langle \text{list1} \rangle} (\pi_{\langle \text{list2} \rangle} (R)) = \pi_{\langle \text{list1} \rangle} (R)$
as long as $\langle \text{list1} \rangle$ contains the attributes from $\langle \text{list2} \rangle$



JOIN operations



JOIN Operation

- Often you require joining two relations.
- Most common form of JOIN is Natural Join.
- Here is one explanation to the JOIN
 - Consider STUDENT and PROGRAM relations in XIT database.
 - We often require JOINING them. Why?
 - Suppose, we want to know name of the program in a student id=101 studies?
 - What we do here is pick tuple of concerned student from student relation, find “matching” tuple from program relation, and
 - Produce a “joined” single tuple by combining both the tuples?



JOIN Operation

- JOIN is a Binary Operator, and has join condition as a parameter.
- General Syntax: $r \bowtie_{\langle \text{join-cond} \rangle} s$
- Here is how it is expressed for JOIN-ing STUDENT and PROGRAM relations

STUDENT $\bowtie_{\langle s.\text{progid}=p.\text{pid} \rangle}$ **PROGRAM**



Example JOIN

StudentID	Name	ProgID	CPI
101	Rahul	BCS	7.5
102	Vikash	BIT	8.6
103	Shally	BEE	5.4
104	Alka	BIT	6.8
105	Ravi	BCS	6.5

- Given STUDENT and PROGRAM relations, following JOIN

student ⋈_{<progid=pid>} **program**
results into following relation

- Given below is resultant relation.

PID	ProgName	Intake	DID
BCS	BTech(CS)	40	CS
BIT	BTech(IT)	30	CS
BEE	BTech(EE)	40	EE
BME	BTech(ME)	40	ME

How do you interpret each tuple in this result relation?

studid character	name character varying(20)	progid character	cpi numerical	pid character	pname character varying	intake smallint	did character
101	Rahul	BCS	8.70	BCS	BTech (CS)	30	CS
102	Vikash	BEC	6.80	BEC	BTech (ECE)	40	EE
103	Shally	BEE	7.40	BEE	BTech (EE)	40	EE
104	Alka	BEC	7.90	BEC	BTech (ECE)	40	EE
105	Ravi	BCS	9.30	BCS	BTech (CS)	30	CS



JOIN operation

- Can be understood through following pseudo algorithm for computing $r \bowtie_{\langle \text{join-cond} \rangle} s$ -

Result set **rs = NULL;**

For each tuple t1 in relation r1

For each tuple t2 in relation r2

If join-cond met then

form new tuple t = t1 + t2

append t to rs



JOIN Operation

- Result of a JOIN has
 - Schema = $R \cup S$ for a JOIN of R and S
 - Number of tuples in the result relation?



JOIN in SQL

- The join $r \bowtie_{\langle \text{join-cond} \rangle} s$ in relational algebra will be written in SQL as following-

SELECT * FROM r JOIN s ON (<join-cond>);

- Example JOIN **student** $\bowtie_{\langle \text{progid}=\text{pid} \rangle}$ **program** is expressed in SQL, as following-

SELECT * FROM STUDENT JOIN PROGRAM ON (progid=pid)



JOIN - Why do we need ?

- Normally database contains *normalized* relations, and in order to answer a query we may need to collect data from multiple relations
- In example on previous slide, we wanted to list program-name, its department-name. But these attributes are in two different relations (Why in two different relations?) for such situations, where we need to combine data from multiple relations, we need joins.



JOIN Operation

- In theory JOIN operation is also explained through CROSS JOIN.
- CROSS JOIN is cross product of operand relations

$$r \bowtie_{\langle \text{join-cond} \rangle} s$$

$$= \sigma_{\langle \text{join-cond} \rangle}(r \times s)$$



CROSS JOIN

student × program

studid character	name character varying(20)	progid character	cpi numerical	pid character	pname character varying	intake smallint	did character
101	Rahul	BCS	8.70	BCS	BTech (CS)	30	CS
102	Vikash	BEC	6.80	BCS	BTech (CS)	30	CS
103	Shally	BEE	7.40	BCS	BTech (CS)	30	CS
104	Alka	BEC	7.90	BCS	BTech (CS)	30	CS
105	Ravi	BCS	9.30	BCS	BTech (CS)	30	CS
101	Rahul	BCS	8.70	BEC	BTech (ECE)	40	EE
102	Vikash	BEC	6.80	BEC	BTech (ECE)	40	EE
103	Shally	BEE	7.40	BEC	BTech (ECE)	40	EE
104	Alka	BEC	7.90	BEC	BTech (ECE)	40	EE
105	Ravi	BCS	9.30	BEC	BTech (ECE)	40	EE
101	Rahul	BCS	8.70	BEE	BTech (EE)	40	EE
102	Vikash	BEC	6.80	BEE	BTech (EE)	40	EE
103	Shally	BEE	7.40	BEE	BTech (EE)	40	EE
104	Alka	BEC	7.90	BEE	BTech (EE)	40	EE
105	Ravi	BCS	9.30	BEE	BTech (EE)	40	EE
101	Rahul	BCS	8.70	BME	BTech (ME)	40	ME
102	Vikash	BEC	6.80	BME	BTech (ME)	40	ME
103	Shally	BEE	7.40	BME	BTech (ME)	40	ME
104	Alka	BEC	7.90	BME	BTech (ME)	40	ME
105	Ravi	BCS	9.30	BME	BTech (ME)	40	ME



Try interpreting content of each tuple of CROSS JOIN? And compare highlighted tuple with tuples in JOIN result

**SELECT * FROM student
CROSS JOIN program**

charact	character varying(20)	character	cpi numeri	pid charac	pname character vary	intake smallin	did chara
101	Rahul	BCS	8.70	BCS	BTech (CS)	30	CS
102	Vikash	BEC	6.80	BCS	BTech (CS)	30	CS
103	Shally	BEE	7.40	BCS	BTech (CS)	30	CS
104	Alka	BEC	7.90	BCS	BTech (CS)	30	CS
105	Ravi	BCS	9.30	BCS	BTech (CS)	30	CS
101	Rahul	BCS	8.70	BEC	BTech (ECE)	40	EE
102	Vikash	BEC	6.80	BEC	BTech (ECE)	40	EE
103	Shally	BEE	7.40	BEC	BTech (ECE)	40	EE
104	Alka	BEC	7.90	BEC	BTech (ECE)	40	EE
105	Ravi	BCS	9.30	BEC	BTech (ECE)	40	EE
101	Rahul	BCS	8.70	BEE	BTech (EE)	40	EE
102	Vikash	BEC	6.80	BEE	BTech (EE)	40	EE
103	Shally	BEE	7.40	BEE	BTech (EE)	40	EE
104	Alka	BEC	7.90	BEE	BTech (EE)	40	EE
105	Ravi	BCS	9.30	BEE	BTech (EE)	40	EE
101	Rahul	BCS	8.70	BME	BTech (ME)	40	ME
102	Vikash	BEC	6.80	BME	BTech (ME)	40	ME
103	Shally	BEE	7.40	BME	BTech (ME)	40	ME
104	Alka	BEC	7.90	BME	BTech (ME)	40	ME
105	Ravi	BCS	9.30	BME	BTech (ME)	40	ME



JOIN and CROSS JOIN (PRODUCT)

- Note that tuples where progid and pid match only represents correct fact set of values.
- It should be easy to observe following:

$\sigma_{\langle \text{progid}=\text{pid} \rangle} (\text{student} \times \text{program})$ is equal to
 $\text{student} \bowtie_{\langle \text{progid}=\text{pid} \rangle} \text{program}$



CROSS JOIN

- Following pseudo algorithm for CROSS PRODUCT or CROSS JOIN

$r \times s$ -

```
result_set = NULL;  
for each tuple t1 in relation r1  
    for each tuple t2 in relation r2  
        form new tuple t = t1 + t2  
        append t to result_set
```

- Below is how CROSS PRODUCT written in SQL-

```
SELECT * FROM student CROSS JOIN program; or  
SELECT * FROM student, program;
```



Do Not perform JOIN through CROSS Product in SQL!

- When you want to perform JOIN in SQL, use JOIN keyword.
- Even though following two are algebraically same. You should be using second style.

```
SELECT * FROM program, department  
WHERE program.did = department.did;
```



```
SELECT * FROM program JOIN department  
ON (program.did = department.did);
```



Some examples



Interpret tuples in results of two JOINS on EMPLOYEE and DEPARTMENT

SELECT * FROM employee AS e JOIN department AS d ON(e.dno=d.dno);

ename character varying(20)	ssn integer	bdate date	gender character(1)	salary numeric(8,2)	superssn integer	dno smallint	dname character varying(20)	dno smallint	mgrssn integer	mgrstartdate date
James	105	1927-06-19	M	55000		1	Headquater	1	105	1971-06-19
Franklin	102	1945-05-22	M	40000	105	5	Research	5	102	1978-05-22
Jennifer	106	1985-01-01	F	43000	105	4	Administration	4	106	1985-01-01
John	101	1955-05-09	M	30000	102	5	Research	5	102	1978-05-22
Alicia	108	1958-04-02	F	25000	106	4	Administration	4	106	1985-01-01
Ramesh	104	1952-08-17	M	38000	102	5	Research	5	102	1978-05-22
Joyce	103	1962-08-17	F	25000	102	5	Research	5	102	1978-05-22
Ahmad	107	1959-09-09	M	25000	106	4	Administration	4	106	1985-01-01

SELECT * FROM employee AS e JOIN department AS d ON(e.ssn=d.mgrssn);

ename character varying(20)	ssn integer	bdate date	gender character(1)	salary numeric(8,2)	superssn integer	dno smallint	dname character varying(20)	dno smallint	mgrssn integer	mgrstartdate date
Franklin	102	1945-05-22	M	40000	105	5	Research	5	102	1978-05-22
Jennifer	106	1985-01-01	F	43000	105	4	Administration	4	106	1985-01-01
James	105	1927-06-19	M	55000		1	Headquater	1	105	1971-06-19



Interpret tuples in result of following JOIN

```
SELECT * FROM employee AS e JOIN employee AS s ON(e.superssn=s.ssn);
```

ename character va	ssn integ	bdate date	gen cha	salary numeric(super integ	dno smal	ename character va	ssn integer	bdate date	gen cha	salary numeric(superssn integer	dno small
Franklin	102	1945-	M	40000	105	5	James	105	1927-	M	55000		1
Jennifer	106	1931-	F	43000	105	4	James	105	1927-	M	55000		1
John	101	1955-	M	30000	102	5	Franklin	102	1945-	M	40000	105	5
Alicia	108	1958-	F	25000	106	4	Jennifer	106	1931-	F	43000	105	4
Ramesh	104	1952-	M	38000	102	5	Franklin	102	1945-	M	40000	105	5
Joyce	103	1962-	F	25000	102	5	Franklin	102	1945-	M	40000	105	5
Ahmad	107	1959-	M	25000	106	4	Jennifer	106	1931-	F	43000	105	4



Interpret tuples in results following JOIN

SELECT * FROM employee AS e JOIN employee AS s ON(e.superssn=s.ssn);

ename character va	ssn integ	bdate date	gen cha	salary numeric(super integ	dno small	ename character va	ssn integer	bdate date	gen cha	salary numeric(superssn integer	dno small
Franklin	102	1945-	M	40000	105	5	James	105	1927-	M	55000		1
Jennifer	106	1931-	F	43000	105	4	James	105	1927-	M	55000		1
John	101	1955-	M	30000	102	5	Franklin	102	1945-	M	40000	105	5
Alicia	108	1958-	F	25000	106	4	Jennifer	106	1931-	F	43000	105	4
Ramesh	104	1952-	M	38000	102	5	Franklin	102	1945-	M	40000	105	5
Joyce	103	1962-	F	25000	102	5	Franklin	102	1945-	M	40000	105	5
Ahmad	107	1959-	M	25000	106	4	Jennifer	106	1931-	F	43000	105	4



Interpret tuples in results of following JOIN

STUDENT ⋈_{<progid=did>} PROGRAM ⋈_{<p.did=d.did>} DEPARTMENT

select * from student JOIN program as p ON (progid=pid)
JOIN department as d ON p.did=d.did;

studid character	name character	progid character	cpi numerical	pid character	pname character varying	intake smallint	did character	did character	dname character varying(30)
101	Rahul	BCS	8.70	BCS	BTech (CS)	30	CS	CS	Computer Engineering
102	Vikash	BEC	6.80	BEC	BTech (ECE)	40	EE	EE	Electrical Engineering
103	Shally	BEE	7.40	BEE	BTech (EE)	40	EE	EE	Electrical Engineering
104	Alka	BEC	7.90	BEC	BTech (ECE)	40	EE	EE	Electrical Engineering
105	Ravi	BCS	9.30	BCS	BTech (CS)	30	CS	CS	Computer Engineering



Other (names/types of) Joins

- Recall that join is expressed as $r \bowtie_{\langle \text{join-cond} \rangle} s$
- Note that Degree of result will be degree of R + degree of S, and cardinality of result can be anything between zero rows in r times rows in s.
- Following are the names/references are found in the literature for JOIN.
- Classically this is general form of JOIN and referred as **Theta Join**.
- If Join condition only includes equality check (that is almost the case) then it is called as **Equi-Join**.



JOIN

- Out of following SQL query shown below

```
SELECT * FROM student AS s JOIN  
program AS p ON (s.progid=p.pid)
```

studid	name	progid	cpi	pid	pname	intake	did
character varying(20)	character varying(20)	character varying(20)	numeric(4,2)	character varying(20)	character varying(20)	smallint	character varying(20)
101	Rahul	BCS	8.70	BCS	BTech (CS)	30	CS
102	Vikash	BEC	6.80	BEC	BTech (ECE)	40	EE
103	Shally	BEE	7.40	BEE	BTech (EE)	40	EE
104	Alka	BEC	7.90	BEC	BTech (ECE)	40	EE
105	Ravi	BCS	9.30	BCS	BTech (CS)	30	CS

- Note the double appearance of column **progid** in join result here.



Natural Join

- Relational model defines NATURAL JOIN that has implicit join condition, that is equality of all common attributes. Below is an SQL example for this.
- This also drops duplicate columns

```
SELECT * FROM program
      NATURAL JOIN department;
```

```
pmjat=# SELECT * FROM program NATURAL JOIN department;
 did | pid |  pname      | intake |      dname
-----+-----+-----+-----+-----
  CS  | BCS | BTech(CS)   |    30  | Computer Engineering
  EE  | BEC | BTech(ECE)  |    40  | Electrical Engineering
  EE  | BEE | BTech(EE)   |    40  | Electrical Engineering
  ME  | BME | BTech(ME)   |    40  | Mechanical Engineering
(4 rows)
```



Natural Join

- Note that SQL Natural Join requires having same name of attributes in both the relations
 - However, in practice, relations may not have same name for FK-PK pairs; for example consider following-
 - progid (in student) and pid (in program), and
 - mgrssn (in department) and ssn (in employee)
- Therefore, we may not always be able to use natural join keyword in SQL
- There are often reasons of choosing different names for semantically same attributes. What?



Natural Join - example

- Following will not give correct result ?

SELECT * FROM student **NATURAL JOIN** program;

- There are no common attributes, therefore it results into a cross product?
- You need to get back to JOIN –

SELECT * FROM student **JOIN** program
ON (student.progid = program.pid);



Operations on relations

- SELECTION: $\sigma_{\langle \text{selection condition} \rangle}(\mathbf{r})$
- PROJECTION: $\pi_{\langle \text{attribute list} \rangle}(\mathbf{r})$
- JOIN: $\mathbf{r} \bowtie_{\langle \text{join-cond} \rangle} \mathbf{s}$
- NATURAL JOIN: $\mathbf{r} * \mathbf{s}$
Requires that \mathbf{r} and \mathbf{s} have a common set of attribute – normally FK-CK pair



Equi (Theta) JOIN and Natural Join

r1	a	b	c
	a1	b1	c1
	a2	(Null)	c2

r2	b	d
	b1	d1
	b2	d2

- `select * from r1 join r2 on (r1.b=r2.b) .`

	a	b	c	b	d
	a1	b1	c1	b1	d1

- `select * from r1 natural join r2`

	b	a	c	d
	b1	a1	c1	d1



Types of JOIN?

- Theta JOIN: can have any boolean expressions in join condition (rarely used)
- EQUI-JOIN: has only equality condition (commonly used)
- Natural JOIN: implicit equality condition on common attributes (and drops duplicate columns)
- INNER JOIN
- OUTER JOIN: LEFT, RIGHT, and FULL



INNER and OUTER JOIN

- Theta Join and Natural Join are inner joins.

- Following result same relation

```
select * from student JOIN program  
ON (progid = pid);
```

```
select * from student INNER JOIN program  
ON (progid = pid);
```



INNER and OUTER JOIN

- INNER JOIN does not join, when
 - Tuples from operand relations do not agree on JOIN-Condition, or
 - Null is found in any of joining attributed
- OUTER join still performs join such cases (in above situations).
- There are three types of OUTER JOIN: LEFT, RIGHT, and FULL based on the way non-matching tuples (and NULLs) are joined and included in the result-set.



Recall- Logical Algo of INNER JOIN

- Iterate through tuples of one relation, and look into other relation for matched tuples, and JOIN wherever there is a match.

```
result_set = NULL;  
for each tuple t1 in relation r1  
  for each tuple t2 in relation r2  
    if join-cond met then  
      form new tuple t = t1 + t2  
      append t to result_set
```



Logical Algo for LEFT OUTER JOIN

- Perform JOIN for all tuples from LEFT operand, whether match or no match.

```
result_set = NULL;
for each tuple t1 in relation r1
    for each tuple t2 in relation r2
        if join-cond met then
            form new tuple t = t1 + t2
            append t to result_set
for each tuple t1 in relation r1
    if t1 NOT IN (result_set)
        form new tuple t = t1 + <null>
        append t to result_set
```



Logical Algo-2 for LEFT OUTER JOIN

- Perform JOIN for all tuples from LEFT operand, whether match or no match.

```
result_set = NULL;
for each tuple t1 in relation r1
    if attribs(left-relation) has NULL then
        form new tuple t = t1 + <null>
        append t to result_set
    match=false
    for each tuple t2 in relation r2
        if join-cond met then
            form new tuple t = t1 + t2
            append t to result_set
            match=true
    if not match then
        form new tuple t = t1 + <null>
        append t to result_set
```

Example INNER JOIN and LEFT JOIN

SELECT * FROM employee AS e [INNER] JOIN employee AS s ON(e.superssn=s.ssn);

ename character va	ssn integ	bdate date	gen cha	salary numeric(super integ	dno smal	ename character va	ssn integer	bdate date	gen cha	salary numeric(superssn integer	dno small
Franklin	102	1945-	M	40000	105	5	James	105	1927-	M	55000		1
Jennifer	106	1931-	F	43000	105	4	James	105	1927-	M	55000		1
John	101	1955-	M	30000	102	5	Franklin	102	1945-	M	40000	105	5
Alicia	108	1958-	F	25000	106	4	Jennifer	106	1931-	F	43000	105	4
Ramesh	104	1952-	M	38000	102	5	Franklin	102	1945-	M	40000	105	5
Joyce	103	1962-	F	25000	102	5	Franklin	102	1945-	M	40000	105	5
Ahmad	107	1959-	M	25000	106	4	Jennifer	106	1931-	F	43000	105	4

SELECT * FROM employee AS e LEFT JOIN employee AS s ON(e.superssn=s.ssn);

ename character va	ssn integ	bdate date	gen cha	salary numeric(super integ	dno smal	ename character va	ssn integer	bdate date	gen cha	salary numeric(superssn integer	dno small
James	105	1927-	M	55000		1							
Franklin	102	1945-	M	40000	105	5	James	105	1927-	M	55000		1
Jennifer	106	1931-	F	43000	105	4	James	105	1927-	M	55000		1
John	101	1955-	M	30000	102	5	Franklin	102	1945-	M	40000	105	5
Alicia	108	1958-	F	25000	106	4	Jennifer	106	1931-	F	43000	105	4
Ramesh	104	1952-	M	38000	102	5	Franklin	102	1945-	M	40000	105	5
Joyce	103	1962-	F	25000	102	5	Franklin	102	1945-	M	40000	105	5
Ahmad	107	1959-	M	25000	106	4	Jennifer	106	1931-	F	43000	105	4



Logical Algo for **RIGHT OUTER JOIN**

- Perform JOIN for all tuples from LEFT operand, whether match or no match.

```
result_set = NULL;
for each tuple t1 in relation r1
    for each tuple t2 in relation r2
        if join-cond met then
            form new tuple t = t1 + t2
            append t to result_set
for each tuple t2 in relation r2
    if t2 NOT IN (result_set)
        form new tuple t = <null> + t2
        append t to result_set
```



Logical Algo of FULL [OUTER] JOIN

- UNION of result of LEFT JOIN and RIGHT JOIN
- Any of the algo used of LEFT or RIGHT can be used, and remaining rows from other side are also included with null values in corresponding attributes



Logical Algo for FULL OUTER JOIN

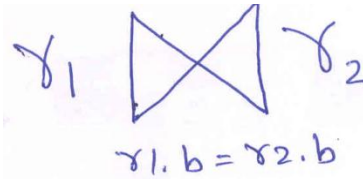
```
result_set = NULL;
for each tuple t1 in relation r1
    for each tuple t2 in relation r2
        if join-cond met then
            form new tuple t = t1 + t2
            append t to result_set
for each tuple t1 in relation r1
    if t1 NOT IN (result_set)
        form new tuple t = t1 + <null>
        append t to result_set
for each tuple t2 in relation r2
    if t2 NOT IN (result_set)
        form new tuple t = <null> + t2
        append t to result_set
```



JOINS: INNER, LEFT, RIGHT, and FULL

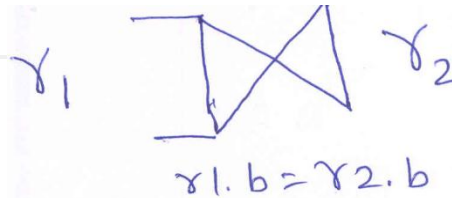
r1		
a	b	c
a1	b1	c1
a2	(Null)	c2

r2	
b	d
b1	d1
b2	d2



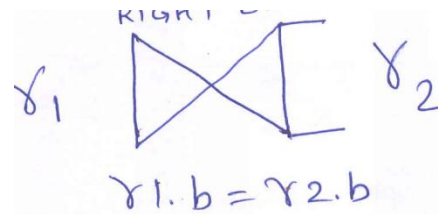
[INNER] JOIN

a	b	c	b1	d
a1	b1	c1	b1	d1



LEFT [OUTER] JOIN

a	b	c	b1	d
a1	b1	c1	b1	d1
a2	(Null)	c2	(Null)	(Null)



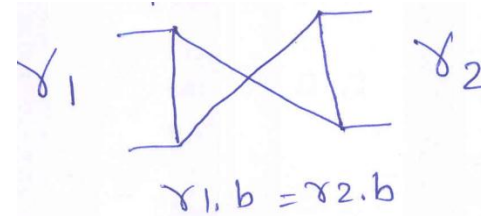
RIGHT [OUTER] JOIN

a	b	c	b1	d
a1	b1	c1	b1	d1
(Null)	(Null)	(Null)	b2	d2



JOINS: INNER, LEFT, RIGHT, and FULL

r1			
	a	b	c
▶	a1	b1	c1
	a2	(Null)	c2



r2		
	b	d
▶	b1	d1
	b2	d2

	a	b	c	b1	d
▶	a1	b1	c1	b1	d1
	a2	(Null)	c2	(Null)	(Null)
	(Null)	(Null)	(Null)	b2	d2

FULL [OUTER] JOIN

FULL JOIN is equivalent to UNION of LEFT and RIGHT JOIN



JOINS in SQL: INNER, LEFT, RIGHT, and FULL

- INNER:

`r1 [INNER] * JOIN r2 on (r1.b=r2.b) ;`

- LEFT OUTER:

`r1 LEFT [OUTER] * JOIN r2 on (r1.b=r2.b) ;`

- RIGHT OUTER:

`r1 RIGHT [OUTER] * JOIN r2 on (r1.b=r2.b) ;`

- FULL OUTER:

`r1 FULL [OUTER] * JOIN r2 on (r1.b=r2.b) ;`

*Optional



Some Exercises!



Ordering of JOINS in FROM clause

- Join operation is *non-associative*. Ordering of evaluation of joins in a FROM clause of SELECT statement is left to right, if there are multiple joins.
- `SELECT ssn, fname, pname AS Project, dname AS "Controlling Dept", hours FROM works_on
NATURAL JOIN project NATURAL JOIN department
JOIN employee ON essn = ssn;`
- Above query mean as below –
`SELECT ssn, fname, pname AS Project, dname AS "Controlling Dept", hours FROM ((works_on
NATURAL JOIN project) NATURAL JOIN
department) JOIN employee ON essn = ssn);`



SET operations



SET Operations

- UNION

- $A \cup B$

Requirement:

UNION Type Compatibility between operands for these operations

- INTERSECT

- $A \cap B$

- EXCEPT (MINUS)

- $A - B$



Type Compatibility for Set operations

- The operand relations $R1(A1, A2, \dots, An)$ and $R2(B1, B2, \dots, Bn)$ must have the same number of attributes, and the domains of corresponding attributes must be compatible; that is, $\text{dom}(Ai) = \text{dom}(Bi)$ for $i=1, 2, \dots, n$.
- The resulting relation for $r1 \cup r2$ has the same attribute names as the *first* operand relation $R1$
- This applies to Intersection and subtraction as well



Example – SET operations

- (a) Two union-compatible relations.
- (b) $\text{STUDENT} \cup \text{INSTRUCTOR}$.
- (c) $\text{STUDENT} \cap \text{INSTRUCTOR}$.
- (d) $\text{STUDENT} - \text{INSTRUCTOR}$
- (e) $\text{INSTRUCTOR} - \text{STUDENT}$

(a)

STUDENT	FN	LN
	Susan	Yao
	Ramesh	Shah
	Johnny	Kohler
	Barbara	Jones
	Amy	Ford
	Jimmy	Wang
	Ernest	Gilbert

INSTRUCTOR	FNAME	LNAME
	John	Smith
	Ricardo	Browne
	Susan	Yao
	Francis	Johnson
	Ramesh	Shah

(b)

FN	LN
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert
John	Smith
Ricardo	Browne
Francis	Johnson

(c)

FN	LN
Susan	Yao
Ramesh	Shah

(d)

FN	LN
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

(e)

FNAME	LNAME
John	Smith
Ricardo	Browne
Francis	Johnson

Courtesy: Elmasri/Navathe



UNION, INTERSECT, MINUS in SQL

- SQL has keywords -
 - UNION for union
 - INTERSECT for intersection
 - EXCEPT for minus
- Syntax:
SELECT . . .
[UNION/INTERSECT/EXCEPT]
SELECT . . . ;



Examples

- Employee that are either manager or supervisor
- Students either study in BCS or BIT
- Employee that are not manager
- Employee that are manager also



NATURAL JOIN and INTERSECTION

- NATURAL JOIN is basically a INTERSECTION problem?
- EMP NATURAL JOIN DEP
==> take INTERSECTION of both sets by checking only dno in both sets and combine the tuples



UNION, INTERSECT, MINUS in SQL

- Because of type compatibility, use of these operations directly have limited use in practice
- In most cases UNION can be performed by having OR in tuple SELECTION criteria (in WHERE Clause of SQL)
- INTERSECT is accomplished by NATURAL JOIN or SEMI JOIN (IN in SQL)
- EXCEPT could be accomplished by SEMI Difference (NOT IN of SQL)
- DISTINCT is implied in SET operations in SQL